

Data Science

BSc Computer Science - Durham University

Matthew Emerson

October 2026

Contents

1	Module Introduction	1
1.1	Assessment	1
1.2	Timetable	1
2	Probability and Statistics	1
2.1	Probability and Statistics: Basic Concepts	1
2.2	Dependency and Bayes Theorem	4
2.3	Discrete Probability Distributions	4
2.4	Continuous Probability Distributions	4
2.5	Gaussian Distribution and its Characteristics	4
2.6	Correlations	4
2.7	Central Limit Theorem and Sampling	4
2.8	Hypothesis Testing	4
2.9	Data Types and Visualisation	4
2.10	Basic Data Analysis with Python	4
3	Data Cleaning and Analytics	4
3.1	Data Quality and How to Produce High-Quality Data	4
3.2	Working with Various File Types	5
3.3	Pandas and Various Features of NumPy	5
3.4	Data Cleaning and Preparation	5
3.5	Python Data Visualisation Libraries	5
3.6	Data Analytics and Reporting	5
4	Computer Graphics	5
5	Image Processing	5
I	Appendix	6
6	Bibliography	6
6.1	Module Professors	6
6.2	References	6
6.3	List of Figures	6
6.4	Reading List	6

1 Module Introduction

Data Science has the module code COMP2271 and is led by Professor Thomas Erlebach.

1.1 Assessment

This module is assessed through two pieces of summative coursework and two final exams, these can be seen in Table 1. Additional formatives are also shown in italics.

Title	Opening	Due	Feedback	Professor	Weight
Data Cleaning & Analytics	Week 11	-	11th Feb	Dr Igor Razgon	25%
Image Processing	15th Jan	2nd Mar	6th May	Dr George Koulieris	25%
Probability and Graphics	May Exam	-	-	Razgon & Koulieris	50%

Table 1: Assessments for Data Science.

Data Cleaning & Analytics Benchtest Information: The Data Cleaning & Analytics test in Week 11 is a closed book, multiple choice test, and will be held during practicals (mine is Friday 4pm).

- Closed-book assessment over one hour. **Check information closer to the time.**
- Comprised of 15 multiple-choice questions, split equally between: basic NumPy, basic Pandas, data cleaning, data analytics, and visualisation.

1.2 Timetable

This module is taught through two one-hour lectures and one two-hour computer class each week as seen below in Table 2.

Day	Time	Class Type	Location	Lecturer
Monday	12:00 - 13:00	Lecture	W103	Dr Igor Razgon
Friday	10:00 - 11:00	Lecture	W103	Dr Igor Razgon
Friday	16:00 - 18:00	Computer Class	CC-0007	-

Table 2: Michaelmas Term Timetable for Data Science.

2 Probability and Statistics

This section is taught by Dr Igor Razgon. The recommended reading for this topic is:

- Probability and Statistics (Degroot, M.H. and Schervish, M.J., 2012).

2.1 Probability and Statistics: Basic Concepts

2.1.1 Counting Theory

Definition 2.1.1.1: The product rule states that, if a procedure can be broken down into a sequence of two tasks with n_1 ways to do the first task and n_2 ways to do the second task, then there are $n_1 n_2$ ways to complete the entire procedure.

Example 2.1.1.1: John wants to fly to Canada from Germany with a connecting flight through the UK. There are 3 flights from Germany to the UK, and 2 flights from the UK to Canada. How many possible flight combinations are there?

Solution: $3 \times 2 = 6$.

□

Definition 2.1.1.2: The sum rule states that, if you consider two independent sub-tasks (we only have to do one), if the first can be done in n_1 ways, and the second can be done in n_2 ways, then the total number of ways to complete the task is $n_1 + n_2$.

Example 2.1.1.2: Mary wants to choose a student project from three lists: lecturer 1 offers 5 projects, lecturer 2 offers 2 projects, and lecturer 3 offers 10 projects. How many possible options for her?

Solution: $5 + 2 + 10 = 17$.

□

Definition 2.1.1.3: The inclusion-exclusion principle states that, if two tasks can be done at the same time (unlike the sum rule in which they are independent), then there is a possibility of overcounting since some ways may be counted twice. To counter this, we add up the number of ways to do each task, then subtract the number of ways to do both tasks.

$$|A \cup B| = (|A| + |B|) - (|A \cap B|) \quad (1)$$

$$|A \cup B \cup C| = (|A| + |B| + |C|) - (|A \cap B| + |A \cap C| + |B \cap C|) + (|A \cap B \cap C|) \quad (2)$$

Definition 2.1.1.4: An r-permutation of a set of r distinct objects is one ordered arrangement. We denote the number of r-permutations of a set with n elements in as $P(n, r)$.

$$P(n, r) = n \times (n - 1) \times (n - 2) \times \dots \times (n - r + 1) = \frac{n!}{(n - r)!} \quad (3)$$

Suppose that a set consists of n objects, of which n_1 are of one type (indistinguishable from one another), n_2 are of a second type, ..., and n_k are of a k_{th} type where $n = n_1 + n_2 + \dots + n_k$. The number of permutations of the objects can be calculated with (4).

$$P(n; n_1, n_2, \dots, n_k) = \frac{n!}{n_1! n_2! \dots n_k!} \quad (4)$$

Definition 2.1.1.5: A r-combination is a counting of distinct objects in an unordered arrangement of length r . The number of r combinations of a set with n distinct elements is given by (5).

$$C(n, r) = \binom{n}{r} = \frac{n!}{(n - r)! \times r!} \quad (5)$$

Example 2.1.1.3: Assume elements A_1, A_2, A_3, A_4 are in a set. We want to choose 2 elements from them.

Solution:

- If the order of elements is important: $P(4, 2) = \frac{4!}{(4-2)!} = 12$.
- If the order of elements is not important: $C(4, 2) = \frac{4!}{(4-2)!2!} = 6$.

□

2.1.2 Probability Theory

Definition 2.1.2.1: Probability is the formalisation of uncertain events. It enables one to infer probabilities of interest based on assumptions and observations.

Definition 2.1.2.2: Statistics is the practice of the science of collecting and analysing data. It deals with experimental data and extracting knowledge from it.

Definition 2.1.2.3: An experiment is a procedure that yields one of a given set of possible outcomes.

Definition 2.1.2.4: The sample space of the experiment is the set of possible outcomes.

Definition 2.1.2.5: An event is any subset of the sample space.

Theorem 2.1.2.1: The Theory of Probability states that, given a finite sample space of all outcomes S , assuming that all outcomes are equally likely, for an event set E where $E \subseteq S$, then the probability of occurrence of the event is given by (6).

$$P(E) = \frac{|E|}{|S|} \quad (6)$$

Proof: Proof omitted. □

Definition 2.1.2.6: The axioms of probability are as follows:

- $P(E)$ is always between 0 and 1: $0 \leq P(E) \leq 1$.
- $P(S) = 1$.
- For disjoint events $E_1, E_2, \dots, E_n$ (with no intersection), $P(E_1 \cup E_2 \cup \dots \cup E_n) = \sum_{i=1}^n P(E_i)$.

Then, from the Axioms of Probability above, we can derive the following:

- **Definition 2.1.2.7:** The complement rule states that $P(\overline{E}) = P(S - E) = P(S) - P(E) = 1 - P(E)$.
- **Definition 2.1.2.8:** The impossible rule states that $P(\emptyset) = 0$.
- **Definition 2.1.2.9:** The general addition rule states that $P(E_1 \cup E_2) = P(E_1) + P(E_2) - P(E_1 \cap E_2)$.
- If $E_1 \subseteq E_2$, then $P(E_1) \leq P(E_2)$.

Example 2.1.2.1: A sequence of 10 bits is randomly generated. What is the probability that at least 1 of these is 0?

Solution: Let E be an event in which all bits are 1.

$$P(E) = \frac{|E|}{|S|} = \frac{1}{2^{10}}$$

Then, by the complement rule, $P(\overline{E}) = P(1 - E)$ where $\overline{E}$ is the event in which at least one bit is not 1.

$$P(\overline{E}) = 1 - \frac{1}{1024} = \frac{1023}{1024}$$

□

2.2 Dependency and Bayes Theorem

2.3 Discrete Probability Distributions

2.4 Continuous Probability Distributions

2.5 Gaussian Distribution and its Characteristics

2.6 Correlations

2.7 Central Limit Theorem and Sampling

2.8 Hypothesis Testing

2.9 Data Types and Visualisation

2.10 Basic Data Analysis with Python

3 Data Cleaning and Analytics

This section is taught by Dr Igor Razgon. The recommended reading for this topic is:

- Python for Data Analytics (McKinney, 2022).

3.1 Data Quality and How to Produce High-Quality Data

3.1.1 Data Quality

Definition 3.1.1.1: Data cleansing is the process of detecting and correcting (or removing) corrupt or inaccurate records from a record set, table, or database and refers to identifying incomplete, inaccurate, or irrelevant parts of the data and then replacing, modifying, or deleting the dirty or coarse data.

For data to be considered high-quality, it should meet the following criteria:

- **Definition 3.1.1.2:** Validity: The degree to which the measures conform to defined business rules or constraints.
- **Definition 3.1.1.3:** Accuracy: The degree of conformity of a measure to a standard or a true value.
- **Definition 3.1.1.4:** Completeness: The degree to which all required measures are known.
- **Definition 3.1.1.5:** Consistency: The degree to which the data is consistent, within the same data set or across multiple data sets.
- **Definition 3.1.1.6:** Uniformity: The degree to which the data is specified using the same unit of measure.

3.1.2 How to Produce High-Quality Data

Producing high-quality data is an iterative process moving through inspection, cleaning, verifying, and reporting.

Definition 3.1.2.1: Inspection: Detect unexpected, incorrect, and inconsistent data.

- **Definition 3.1.2.2:** Data profiling is the process of viewing the summary statistics of your data to identify standards or patterns in columns of data. Here it is also important to note the data type, any missing values, the number of unique values, the distribution of values, and the relationship between the data and other columns.
- Visualising the data can also be helpful for identifying outliers, as can using statistical methods such as mean, standard deviation, range, and/or quartiles.

Definition 3.1.2.3: Cleaning: Fix or remove the anomalies detected.

- It is important to identify and remove duplicate entries in your data sets.
- When converting the data into a consistent data type, values unable to fit the format should be flagged and converted to NA value.
- Beyond the data types, the data should be further standardised. This means strings are of a consistent capitalisation, numerical values use the same measurement unit, and dates are in the same format (pick either the UK or US format).
- It can also be useful to scale/transform your data to fit within a certain scale (i.e. 0-100 or 0-1). This can make data easier to plot (e.g., exponential data being logged).
- You can either replace missing values by using a mask to indicate where they are missing, or choosing a sentinel value to replace the missing entries. Conversely, you can also drop the whole record from the data set, or where possible calculate the missing values from other observations.
- **Definition 3.1.2.4:** We define outliers as values which fall 1.5 times the interquartile range from both the first and third quartiles.

Definition 3.1.2.5: Verifying: After cleaning, the results are inspected to verify correctness.

Definition 3.1.2.6: Reporting: A report about the changes made and the quality of the currently stored data is recorded.

3.2 Working with Various File Types

3.3 Pandas and Various Features of NumPy

3.4 Data Cleaning and Preparation

3.5 Python Data Visualisation Libraries

3.6 Data Analytics and Reporting

4 Computer Graphics

This section is taught by Dr George Alex Koulieris. [Check for recommended reading.](#)

5 Image Processing

This section is taught by Dr George Alex Koulieris. [Check for recommended reading.](#)

Part I

Appendix

6 Bibliography

6.1 Module Professors

Module Lead: Professor Thomas Erlebach

- Focus: Algorithms and networks.
- Email: thomas.erlebach@durham.ac.uk
- Office Hours: n/a.

Lecturer: Doctor Igor Razgon

- Focus: Algorithms and complexity, graph theory, and discrete mathematics.
- Email: igor.razgon@durham.ac.uk
- Office Hours: Friday 14:00 - 15:00.

Lecturer: Doctor George Alex Koulieris

- Focus: Image processing and computer graphics.
- Email: georgios.a.koulieris@durham.ac.uk
- Office Hours: Please email to make an appointment.

6.2 References

6.3 List of Figures

6.4 Reading List

6.4.1 Probability

Degroot, M.H. and Schervish, M.J. (2012). Probability and Statistics. Boston Mass.: Pearson.

6.4.2 Data Cleaning and Analytics

McKinney, W. (2022). Python for Data Analysis. O'Reilly Media, Inc.