

Year Two Lecture Notes

BSc Computer Science - Durham University

Matthew Emerson

October 2026

Contents

I	Artificial Intelligence	1
1	Module Introduction	1
1.1	Assessment	1
1.2	Timetable	1
2	Machine Learning	1
2.1	Introduction to Machine Learning	1
2.2	Machine Learning Workflow and Coding Environment	3
2.3	Data Understanding	3
2.4	Data Preparation Workflow	3
3	Artificial Intelligence Search	3
3.1	Coding Real-World Problems for Search Tree Expansion	3
3.2	Uninformed Search: BFS, DFS, and Iterative Deepening	3
3.3	General Heuristic Search: Greedy Best-First Search and A* Search	3
3.4	A* Search and Deriving Iterative Deepening A* Search	3
3.5	Heuristic Functions and Their Development	3
3.6	Local Search: Hill-Climbing, Simulated Annealing, and Genetic Algorithms	3
3.7	Swarm Behaviour, How Ants Forage, and a Simple S-ACO algorithm	3
3.8	Applying Different Ant Systems to the TSP	3
3.9	General Versions of Particle Swarm Optimisation and Differential Evolution	3
3.10	Applying Particle Swarm Optimization to the TSP	3
4	Ethics and Bias in AI	3
4.1	Introduction to AI Ethics and Related Concepts	4
4.2	Societal Impacts of AI	4
4.3	Defining Bias and Fairness in ML	4
4.4	How to Build Better Systems	4
4.5	Value Sensitive Design	4
4.6	Evaluations and Audits	4
4.7	Bias Mitigation Strategies	4
4.8	Transparency and Accountability for Trustworthy Systems	4
4.9	Explainable AI	4
4.10	Future Trends and Ethical Frontiers	4
II	Data Science	5
5	Module Introduction	5
5.1	Assessment	5
5.2	Timetable	5
6	Probability and Statistics	5
6.1	Probability and Statistics: Basic Concepts	5
6.2	Dependency and Bayes Theorem	8
6.3	Discrete Probability Distributions	8
6.4	Continuous Probability Distributions	8
6.5	Gaussian Distribution and its Characteristics	8
6.6	Correlations	8
6.7	Central Limit Theorem and Sampling	8
6.8	Hypothesis Testing	8

6.9	Data Types and Visualisation	8
6.10	Basic Data Analysis with Python	8
7	Data Cleaning and Analytics	8
7.1	Data Quality and How to Produce High-Quality Data	8
7.2	Working with Various File Types	9
7.3	Pandas and Various Features of NumPy	9
7.4	Data Cleaning and Preparation	9
7.5	Python Data Visualisation Libraries	9
7.6	Data Analytics and Reporting	9
8	Computer Graphics	9
9	Image Processing	9
III	Networks and Systems	10
10	Module Introduction	10
10.1	Assessment	10
10.2	Timetable	10
11	Cyber Security	10
11.1	Introduction, Threat Modelling, and Applied Cryptography	10
12	Networks	13
13	Databases	13
13.1	Enhanced Entity-Relationship Model	13
13.2	Semistructured Databases (XML)	13
13.3	XML Data Manipulation (XPath and XQuery)	13
13.4	Transactions	13
13.5	Concurrency Control	13
13.6	Distributed Transactions	13
13.7	Distributed Concurrency Control	13
14	Distributed Systems	13
14.1	Introduction to Distributed Systems	13
14.2	Middleware and RMI Technologies	13
14.3	Message-Oriented Middleware and Web Services	13
14.4	Replication Models	13
14.5	Concurrent Events and Consistency Control	13
14.6	Fault Tolerance	13
14.7	Load Distribution	13
IV	Programming Paradigms	14
15	Module Introduction	14
15.1	Assessment	14
15.2	Timetable	14

16 Systems Programming	14
16.1 Version Control and UNIX	14
16.2 The Basics of C	14
16.3 Control Flow and Functions	16
16.4 Data Types and Structures	16
16.5 Advanced Bash Scripting	16
16.6 Pointers	16
16.7 Memory Management	16
16.8 Debugging	16
16.9 Scope and Recursion	16
16.10 Libraries	16
16.11 stdlib	16
16.12 C++	16
17 Functional Programming	16
18 Object Oriented Programming	16
18.1 Java Fundamentals	16
18.2 OOP with Classes and Objects	16
18.3 Encapsulation	16
18.4 Inheritance	16
18.5 Abstraction with Abstract Classes	16
18.6 Abstraction with Interfaces	16
18.7 Polymorphism	16
18.8 Java Collections	16
18.9 Best Practises	16
V Software Engineering	17
19 Module Introduction	17
19.1 Assessment	17
19.2 Timetable	17
20 Project Management	17
21 Software Engineering	17
22 Legal, Social, Ethical, and Professional Issues	17
22.1 Introduction to Software Engineering and Software Engineering Ethics	17
23 Group Project	18
VI Theory of Computation	19
24 Module Introduction	19
24.1 Assessment	19
24.2 Timetable	19
25 Models of Computation	19
25.1 Finite-State Automata and Regular Languages	19
25.2 Büchi Automata and Omega-Regular Languages	21
25.3 Linear-Time Logic and Translating Into Büchi Automata	21

26 Algorithms and Complexity	21
26.1 Asymptotic Complexity and Notation	21
26.2 Greedy Algorithms	21
26.3 Divide and Conquer Algorithms	21
26.4 Dynamic Programming Algorithms	21
26.5 Machinings in Graphs	21
 VII Appendix	 22
27 Artificial Intelligence	22
27.1 Module Professors	22
27.2 References	22
27.3 List of Figures	22
27.4 Reading List	22
 28 Data Science	 23
28.1 Module Professors	23
28.2 References	23
28.3 List of Figures	23
28.4 Reading List	23
 29 Networks and Systems	 23
29.1 Module Professors	23
29.2 References	24
29.3 List of Figures	24
29.4 Reading List	24
 30 Programming Paradigms	 24
30.1 Module Professors	24
30.2 References	25
30.3 List of Figures	25
 31 Software Engineering	 25
31.1 Module Professors	25
31.2 References	26
31.3 List of Figures	26
31.4 Reading List	26
 32 Theory of Computation	 26
32.1 Module Professors	26
32.2 References	27
32.3 List of Figures	27
32.4 Recommended Reading	27

Part I

Artificial Intelligence

1 Module Introduction

Artificial Intelligence has the module code COMP2261, and is led by Dr Yang Long.

1.1 Assessment

Title	Opening	Due	Feedback	Professor	Weight
Machine Learning	9th Oct	15th Dec	21st Jan	Dr Yang Long	50%
AI Search	15th Jan	18th Feb	18th Mar	Prof Iain Stewart	25%
Bias in AI	5th Feb	30th Mar	6th May	Dr Shauna Concannon	25%

Table 1: Assessments for Artificial Intelligence.

1.2 Timetable

This module is taught through two one-hour lectures and one two-hour computer class each week as seen below in Table 2.

Day	Time	Class Type	Location	Lecturer
Monday	14:00 - 16:00	Computer Class	TLC025	-
Tuesday	16:00 - 17:00	Lecture	W103	Dr Yang Long
Friday	15:00 - 16:00	Lecture	W103	Dr Yang Long

Table 2: Michaelmas Term Timetable for Artificial Intelligence.

2 Machine Learning

This section is taught by Dr Yang Long. The recommended reading for this topic is:

- Mathematics for Machine Learning (Deisenroth, M.P., Faisal, A.A., and Ong, C.S., 2020).
- Mastering Machine Learning with scikit-learn (Hackeling, G., 2017).

2.1 Introduction to Machine Learning

2.1.1 Module Information

This module is organised around a direct connection between teaching and practical work. The lectures introduce the concepts, mathematics, assumptions, and experimental consequences needed for that week's practical. The practical then produces the evidence used in the corresponding coursework report. This can be seen in Table 3.

Stage	What you do	What you keep
Lecture	Understand the main question, mechanisms, assumptions, and key equations.	Notes in your own words and questions that still need an answer.
Practical	Implement, test, compare, visualise, and discuss the method.	A notebook that runs from top to bottom with outputs included.
Weekly Report	Select the strongest evidence and explain why the result occurred.	A concise report that links theory, method, result, and interpretation.
Portfolio	Revise the nine reports and connect ideas across the module.	A coherent record of technical progress and developing judgement.

Table 3: The learning and coursework model.

For more information about expectations for evidence, see (Long and Durham University, 2026, pp.3-4).

2.1.2 What is Machine Learning?

Definition 2.1.2.1: Arthur Samuel said in 1959 that machine learning is the

“field of study that gives computers the ability to learn without being explicitly programmed.”

The field arose from the question – can a computer learn on its own in order to perform a specific task, rather than relying on a programmer explicitly setting up rules for it? Similarly, Tom Mitchell defined machine learning in 1997 as

“A computer program is said to learn from experience E with respect to some class of tasks T and performance measure P , if its performance at tasks in T , as measured by P , improves with experience E .”

From this, we can define a well-defined learning task as $\langle T, P, E \rangle$.

There are three typical types of machine learning, these are:

- **Definition 2.1.2.2:** Supervised learning seeks to learn the mapping (rules) between inputs and outputs. Labelled data (correct answers) is provided in past input and output pairs during the learning process to train the model on how it should behave for previously unseen data.
- **Definition 2.1.2.3:** Unsupervised learning seeks to discover the hidden pattern from a set of inputs. Unlabelled data is provided as the past input during the learning process. The program will then attempt to group similar instances into groups.
- **Definition 2.1.2.4:** Reinforcement learning works by providing occasional positive and negative feedback to reinforce behaviours. Good behaviours are rewarded with treats and therefore become more common whereas bad behaviours are punished and therefore become less common. This is a balance between exploration (of uncharted territory) and exploitation (of current knowledge).

2.1.3 Differences between AI, ML, and DL

Definition 2.1.3.1: Artificial intelligence is defined as human intelligence being exhibited by machines.

Definition 2.1.3.2: In simple terms, machine learning can be defined as an approach to achieve artificial intelligence.

Definition 2.1.3.3: Deep learning is defined as a technique for achieving machine learning. As such, we can consider AI to be a subset of ML and ML to be a subset of DL.

2.2 Machine Learning Workflow and Coding Environment

2.2.1 Machine Learning Workflow

Definition 2.2.1.1: To train a model, you use data to create a model and fine-tune it to make better predictions. We can then use this predictive model to make predictions on previously unseen data in order to perform specific tasks.

Add notes while in lecture, because the PPT is useless.

2.2.2 Coding Environment

Add notes while in lecture, because the PPT is useless.

2.3 Data Understanding

2.4 Data Preparation Workflow

3 Artificial Intelligence Search

This section is taught by Professor Iain Stewart. There is no recommended reading for this topic however there may be an additional video to watch after lectures to enable more content to be covered.

Watch the FAQs video.

3.1 Coding Real-World Problems for Search Tree Expansion

Watch ahead of time - this will not be covered in lectures. Week 5 covers no content.

3.2 Uninformed Search: BFS, DFS, and Iterative Deepening

Watch ahead of time - this will not be covered in lectures. Week 5 covers no content.

3.3 General Heuristic Search: Greedy Best-First Search and A* Search

3.4 A* Search and Deriving Iterative Deepening A* Search

3.5 Heuristic Functions and Their Development

3.6 Local Search: Hill-Climbing, Simulated Annealing, and Genetic Algorithms

3.7 Swarm Behaviour, How Ants Forage, and a Simple S-ACO algorithm

3.8 Applying Different Ant Systems to the TSP

3.9 General Versions of Particle Swarm Optimisation and Differential Evolution

3.10 Applying Particle Swarm Optimization to the TSP

4 Ethics and Bias in AI

This section is taught by Dr Shauna Concannon. The recommended reading is handled on a weekly basis for this topic.

- 4.1 Introduction to AI Ethics and Related Concepts
- 4.2 Societal Impacts of AI
- 4.3 Defining Bias and Fairness in ML
- 4.4 How to Build Better Systems
- 4.5 Value Sensitive Design
- 4.6 Evaluations and Audits
- 4.7 Bias Mitigation Strategies
- 4.8 Transparency and Accountability for Trustworthy Systems
- 4.9 Explainable AI
- 4.10 Future Trends and Ethical Frontiers

Part II

Data Science

5 Module Introduction

Data Science has the module code COMP2271 and is led by Professor Thomas Erlebach.

5.1 Assessment

This module is assessed through two pieces of summative coursework and two final exams, these can be seen in Table 4. Additional formatives are also shown in *italics*.

Title	Opening	Due	Feedback	Professor	Weight
Data Cleaning & Analytics	Week 11	-	11th Feb	Dr Igor Razgon	25%
Image Processing	15th Jan	2nd Mar	6th May	Dr George Koulieris	25%
Probability and Graphics	May Exam	-	-	Razgon & Koulieris	50%

Table 4: Assessments for Data Science.

Data Cleaning & Analytics Benchtest Information: The Data Cleaning & Analytics test in Week 11 is a closed book, multiple choice test, and will be held during practicals (mine is Friday 4pm).

- Closed-book assessment over one hour. **Check information closer to the time.**
- Comprised of 15 multiple-choice questions, split equally between: basic NumPy, basic Pandas, data cleaning, data analytics, and visualisation.

5.2 Timetable

This module is taught through two one-hour lectures and one two-hour computer class each week as seen below in Table 5.

Day	Time	Class Type	Location	Lecturer
Monday	12:00 - 13:00	Lecture	W103	Dr Igor Razgon
Friday	10:00 - 11:00	Lecture	W103	Dr Igor Razgon
Friday	16:00 - 18:00	Computer Class	CC-0007	-

Table 5: Michaelmas Term Timetable for Data Science.

6 Probability and Statistics

This section is taught by Dr Igor Razgon. The recommended reading for this topic is:

- Probability and Statistics (Degroot, M.H. and Schervish, M.J., 2012).

6.1 Probability and Statistics: Basic Concepts

6.1.1 Counting Theory

Definition 6.1.1.1: The product rule states that, if a procedure can be broken down into a sequence of two tasks with n_1 ways to do the first task and n_2 ways to do the second task, then there are $n_1 n_2$ ways to complete the entire procedure.

Example 6.1.1.1: John wants to fly to Canada from Germany with a connecting flight through the UK. There are 3 flights from Germany to the UK, and 2 flights from the UK to Canada. How many possible flight combinations are there?

Solution: $3 \times 2 = 6$.

□

Definition 6.1.1.2: The sum rule states that, if you consider two independent sub-tasks (we only have to do one), if the first can be done in n_1 ways, and the second can be done in n_2 ways, then the total number of ways to complete the task is $n_1 + n_2$.

Example 6.1.1.2: Mary wants to choose a student project from three lists: lecturer 1 offers 5 projects, lecturer 2 offers 2 projects, and lecturer 3 offers 10 projects. How many possible options for her?

Solution: $5 + 2 + 10 = 17$.

□

Definition 6.1.1.3: The inclusion-exclusion principle states that, if two tasks can be done at the same time (unlike the sum rule in which they are independent), then there is a possibility of overcounting since some ways may be counted twice. To counter this, we add up the number of ways to do each task, then subtract the number of ways to do both tasks.

$$|A \cup B| = (|A| + |B|) - (|A \cap B|) \quad (1)$$

$$|A \cup B \cup C| = (|A| + |B| + |C|) - (|A \cap B| + |A \cap C| + |B \cap C|) + (|A \cap B \cap C|) \quad (2)$$

Definition 6.1.1.4: An r-permutation of a set of r distinct objects is one ordered arrangement. We denote the number of r-permutations of a set with n elements in as $P(n, r)$.

$$P(n, r) = n \times (n - 1) \times (n - 2) \times \dots \times (n - r + 1) = \frac{n!}{(n - r)!} \quad (3)$$

Suppose that a set consists of n objects, of which n_1 are of one type (indistinguishable from one another), n_2 are of a second type, ..., and n_k are of a k_{th} type where $n = n_1 + n_2 + \dots + n_k$. The number of permutations of the objects can be calculated with (4).

$$P(n; n_1, n_2, \dots, n_k) = \frac{n!}{n_1! n_2! \dots n_k!} \quad (4)$$

Definition 6.1.1.5: A r-combination is a counting of distinct objects in an unordered arrangement of length r . The number of r combinations of a set with n distinct elements is given by (5).

$$C(n, r) = \binom{n}{r} = \frac{n!}{(n - r)! \times r!} \quad (5)$$

Example 6.1.1.3: Assume elements A_1, A_2, A_3, A_4 are in a set. We want to choose 2 elements from them.

Solution:

- If the order of elements is important: $P(4, 2) = \frac{4!}{(4-2)!} = 12$.
- If the order of elements is not important: $C(4, 2) = \frac{4!}{(4-2)!2!} = 6$.

□

6.1.2 Probability Theory

Definition 6.1.2.1: Probability is the formalisation of uncertain events. It enables one to infer probabilities of interest based on assumptions and observations.

Definition 6.1.2.2: Statistics is the practice of the science of collecting and analysing data. It deals with experimental data and extracting knowledge from it.

Definition 6.1.2.3: An experiment is a procedure that yields one of a given set of possible outcomes.

Definition 6.1.2.4: The sample space of the experiment is the set of possible outcomes.

Definition 6.1.2.5: An event is any subset of the sample space.

Theorem 6.1.2.1: The Theory of Probability states that, given a finite sample space of all outcomes S , assuming that all outcomes are equally likely, for an event set E where $E \subseteq S$, then the probability of occurrence of the event is given by (6).

$$P(E) = \frac{|E|}{|S|} \quad (6)$$

Proof: Proof omitted.

□

Definition 6.1.2.6: The axioms of probability are as follows:

- $P(E)$ is always between 0 and 1: $0 \leq P(E) \leq 1$.
- $P(S) = 1$.
- For disjoint events $E_1, E_2, \dots, E_n$ (with no intersection), $P(E_1 \cup E_2 \cup \dots \cup E_n) = \sum_{i=1}^n P(E_i)$.

Then, from the Axioms of Probability above, we can derive the following:

- **Definition 6.1.2.7:** The complement rule states that $P(\overline{E}) = P(S - E) = P(S) - P(E) = 1 - P(E)$.
- **Definition 6.1.2.8:** The impossible rule states that $P(\emptyset) = 0$.
- **Definition 6.1.2.9:** The general addition rule states that $P(E_1 \cup E_2) = P(E_1) + P(E_2) - P(E_1 \cap E_2)$.
- If $E_1 \subseteq E_2$, then $P(E_1) \leq P(E_2)$.

Example 6.1.2.1: A sequence of 10 bits is randomly generated. What is the probability that at least 1 of these is 0?

Solution: Let E be an event in which all bits are 1.

$$P(E) = \frac{|E|}{|S|} = \frac{1}{2^{10}}$$

Then, by the complement rule, $P(\overline{E}) = P(1 - E)$ where $\overline{E}$ is the event in which at least one bit is not 1.

$$P(\overline{E}) = 1 - \frac{1}{1024} = \frac{1023}{1024}$$

□

6.2 Dependency and Bayes Theorem

6.3 Discrete Probability Distributions

6.4 Continuous Probability Distributions

6.5 Gaussian Distribution and its Characteristics

6.6 Correlations

6.7 Central Limit Theorem and Sampling

6.8 Hypothesis Testing

6.9 Data Types and Visualisation

6.10 Basic Data Analysis with Python

7 Data Cleaning and Analytics

This section is taught by Dr Igor Razgon. The recommended reading for this topic is:

- Python for Data Analytics (McKinney, 2022).

7.1 Data Quality and How to Produce High-Quality Data

7.1.1 Data Quality

Definition 7.1.1.1: Data cleansing is the process of detecting and correcting (or removing) corrupt or inaccurate records from a record set, table, or database and refers to identifying incomplete, inaccurate, or irrelevant parts of the data and then replacing, modifying, or deleting the dirty or coarse data.

For data to be considered high-quality, it should meet the following criteria:

- **Definition 7.1.1.2:** Validity: The degree to which the measures conform to defined business rules or constraints.
- **Definition 7.1.1.3:** Accuracy: The degree of conformity of a measure to a standard or a true value.
- **Definition 7.1.1.4:** Completeness: The degree to which all required measures are known.
- **Definition 7.1.1.5:** Consistency: The degree to which the data is consistent, within the same data set or across multiple data sets.
- **Definition 7.1.1.6:** Uniformity: The degree to which the data is specified using the same unit of measure.

7.1.2 How to Produce High-Quality Data

Producing high-quality data is an iterative process moving through inspection, cleaning, verifying, and reporting.

Definition 7.1.2.1: Inspection: Detect unexpected, incorrect, and inconsistent data.

- **Definition 7.1.2.2:** Data profiling is the process of viewing the summary statistics of your data to identify standards or patterns in columns of data. Here it is also important to note the data type, any missing values, the number of unique values, the distribution of values, and the relationship between the data and other columns.
- Visualising the data can also be helpful for identifying outliers, as can using statistical methods such as mean, standard deviation, range, and/or quartiles.

Definition 7.1.2.3: Cleaning: Fix or remove the anomalies detected.

- It is important to identify and remove duplicate entries in your data sets.
- When converting the data into a consistent data type, values unable to fit the format should be flagged and converted to NA value.
- Beyond the data types, the data should be further standardised. This means strings are of a consistent capitalisation, numerical values use the same measurement unit, and dates are in the same format (pick either the UK or US format).
- It can also be useful to scale/transform your data to fit within a certain scale (i.e. 0-100 or 0-1). This can make data easier to plot (e.g., exponential data being logged).
- You can either replace missing values by using a mask to indicate where they are missing, or choosing a sentinel value to replace the missing entries. Conversely, you can also drop the whole record from the data set, or where possible calculate the missing values from other observations.
- **Definition 7.1.2.4:** We define outliers as values which fall 1.5 times the interquartile range from both the first and third quartiles.

Definition 7.1.2.5: Verifying: After cleaning, the results are inspected to verify correctness.

Definition 7.1.2.6: Reporting: A report about the changes made and the quality of the currently stored data is recorded.

7.2 Working with Various File Types

7.3 Pandas and Various Features of NumPy

7.4 Data Cleaning and Preparation

7.5 Python Data Visualisation Libraries

7.6 Data Analytics and Reporting

8 Computer Graphics

This section is taught by Dr George Alex Koulieris. [Check for recommended reading.](#)

9 Image Processing

This section is taught by Dr George Alex Koulieris. [Check for recommended reading.](#)

Part III

Networks and Systems

10 Module Introduction

Networks and Systems has the module code COMP2211 and is led by Dr Wanqing Tu.

10.1 Assessment

This module is assessed through two pieces of summative coursework and two final exams, these can be seen in Table 6. Additional formatives are also shown in *italics*.

Title	Opening	Due	Feedback	Professor	Weight
Cyber Security	9th Oct	5th Nov	3rd Dec	Dr Carlton Shepherd	25%
Networks	23rd Oct	21st Jan	18th Feb	Dr Wanqing Tu	25%
Dist. Sys. & Databases	May Exam	-	-	Dr Vasantam	50%

Table 6: Assessments for Networks and Systems.

10.2 Timetable

This module is taught through one two-hour lecture and one two-hour computer class each week as seen below in Table 7.

Day	Time	Class Type	Location	Lecturer
Thursday	09:00 - 11:00	Lecture	CLC202	Dr Carlton Shepherd
Thursday	16:00 - 18:00	Computer Class	CC-0007	-

Table 7: Michaelmas Term Timetable for Networks and Systems.

11 Cyber Security

This section is taught by Dr Carlton Shepherd. The recommended reading for this module is as follows:

- Practical Cryptography for Developers (Nakov, Stefanov and Shideroff, 2018), specifically the section on “Encryption: Symmetric and Asymmetric”.

11.1 Introduction, Threat Modelling, and Applied Cryptography

11.1.1 Basic Definitions

Definition 11.1.1.1: A system’s security depends on the definition of the adversary: who they are, what they can do (or their capabilities) and their determination (time, money, resources). This is the adversary or threat model.

Definition 11.1.1.2: An asset is something that is worth protecting.

Definition 11.1.1.3: A threat is a person or program that may try and access/steal your assets.

Definition 11.1.1.4: An assumption is something you are relying on to be true/secure by default.

Definition 11.1.1.5: Your goal should be in terms of what needs to be secure, and how long for.

Definition 11.1.1.6: The controls are the things you can do to attempt to protect your assets.

The CIA triad are the three main security properties, these are:

- **Definition 11.1.1.7:** Confidentiality: The outcomes of different actions in the system are only visible by authorised subjects.
- **Definition 11.1.1.8:** Integrity: Actions modifying the system can only be performed by authorised subjects.
- **Definition 11.1.1.9:** Availability: Subjects can perform all actions they are authorised for.

In addition to these properties, it is also important to consider additional properties, such as:

- **Definition 11.1.1.10:** Authenticity: An entity is who they claim they claim to be.
- **Definition 11.1.1.11:** Non-Repudiation: An actor cannot easily deny that an action has occurred.
- **Definition 11.1.1.12:** Context-Binding: An action is valid only for a particular period of time, distance, location, and/or environment.

Definition 11.1.1.13: An exploit is a piece of code, a tool, or a technique that takes advantage of a flaw in software, hardware, or a network to cause unintended behaviour.

Definition 11.1.1.14: A vulnerability is a weakness or flaw in a system's design, implementation, operation, or human behaviour that an attacker can exploit to compromise security.

Definition 11.1.1.15: Risk is the probability or likelihood that a threat will exploit a vulnerability and cause harm, loss, or negative consequences to an asset.

11.1.2 STRIDE Modelling

Definition 11.1.2.1: One method of classifying security risks is STRIDE, which was first developed by Microsoft. It is separated as follows:

- **Definition 11.1.2.1:** Spoofing is the act of using someone else's credentials to perform an operation and is a failure of identification and authentication.
- **Definition 11.1.2.2:** Tampering is the malicious modification of data and is a failure of integrity and authenticity.
- **Definition 11.1.2.3:** Repudiation is the act of denying having performed an operation and is a failure of accountability and identification.
- **Definition 11.1.2.4:** Information Disclosure is the exposure of information to unauthorised users and is a failure of confidentiality.
- **Definition 11.1.2.5:** Denial of Service is the act of denying service to legitimate users and is a failure of availability.
- **Definition 11.1.2.6:** Elevation of Privilege is an instance of an unprivileged user gaining privileged access and is a failure of confinement and authentication.

This is a helpful model for structuring the analysis of potential system threats but it is not exhaustive or exclusive.

11.1.3 Applied Cryptography

Definition 11.1.3.1: To encode some given data, it is transformed from one format/domain to another. This is fully reversible, and does not provide any secrecy. The inverse procedure is called decoding.

- **Definition 11.1.3.2:** A subtype of encoding is steganography, where the existence of the secret message is concealed.

Definition 11.1.3.3: To hash some given data, it is transformed into a (usually) smaller, fixed-size. This is most commonly considered a one-way function. An example of this are SHA algorithms.

- A hashing algorithm works by taking a message m and applying a hashing algorithm $h = H(m)$. h can then be sent securely and the recipient can then take $h' = H(m')$ where $m = m'$, and if $h' = h$, then the message received is correct.

Definition 11.1.3.4: Non-secret data (plaintext) can be encrypted into ciphertext using a secret key. The output is $Enc(m, K) = c$. The inverse procedure, decryption, again uses a secret key to transform the ciphertext back to plaintext $Dec(c, K) = m$.

11.1.4 Cipher Examples

Definition 11.1.4.1: A substitution cipher works by taking each letter in the alphabet and replacing it with another one. This was purportedly created and used by Julius Caesar. These are weak though and can be broken by:

- **Definition 11.1.4.1.1:** A brute-force attack tries every possible key, which for substitution ciphers becomes trivial.
- **Definition 11.1.4.2:** A frequency analysis attack matches the ciphertext to the known frequency of English letters.

Definition 11.1.4.4: Another historical example of a cipher is a Vigenere cipher. This uses a sliding password, such that repeats of the same letter are coded differently. Attacks against this have to guess possible key lengths and then try all possibilities using frequency analysis. This is much harder to break.

Definition 11.1.4.5: Another encryption method is using one time pads. Here, you select a key as long as your message. Data is converted into binary and bitwise XORed with the key. To decode it, you then bitwise XOR the ciphertext with the key. This is considered theoretically secure as long as the key generation process is truly random, the key is never reused, and the key is kept secret.

Definition 11.1.4.6: Stream ciphers use a short secret key and a nonce - a use-once but public value - to generate a long pseudorandom sequence called a keystream. Only the short secret key must be kept private. We then XOR with the message and keystream.

- To encrypt, you take the secret key K and the nonce N and feed them through a keystream generator. The resulting keystream S is then XORed with the plaintext p to produce ciphertext C .
- To decrypt, the recipient takes the same secret key K and nonce N , feeds them through the keystream generator and then XOR the result with the ciphertext C to produce the original plaintext p .

Definition 11.1.4.7: Block ciphers transform a fixed-size block of plaintext into a ciphertext block of the same size, using a secret key. Encryption applies a keyed transformation to the plaintext block, and decryption reverses the transformation using the same key. An example of this is AES encryption. To encrypt longer messages, we use a mode of operation which specifies how to encrypt message blocks in sequence - this usually involves a nonce (initialisation vector).

- **Definition 11.1.4.7:** AES Electronic Code Book (ECB) mode works by taking the plaintext, splitting it into blocks, and then encrypting each of them separately. This is not secure.
- **Definition 11.1.4.8:** AES Cipher Block Chaining (CBC) mode splits the plaintext into blocks and encrypts the first one. The result is then XORed with the second plaintext block before it is encrypted, and so on.

11.1.5 Public Key Cryptography

Definition 11.1.5.1: Rivest, Shamir, Adelman (RSA) is an algorithm where person A computes the product of two prime numbers $n = p \times q$ and makes the value of n public. Person A then generates a public value e and a private value d using p and q . e can then be used to encrypt messages to send to person A, and d used to decrypt them. This is secure based on the computational difficulty of factorisation, but with the rise of quantum computing, is becoming less secure.

Definition 11.1.5.2: Digital signatures enables others to check that a message was signed using a private key and has not been altered. To sign a message, the sender uses their private key, and then sends the message and signature to the recipient. To verify the signature, the recipient uses the sender's public key and compares the result to the message.

12 Networks

This section is taught by Dr Wanqing Tu. [Check for recommended reading.](#)

13 Databases

This section is taught by Dr Vasantam. The recommended reading for this topic is:

- Connolly, T.M. and Begg, C.E. (2015). Database Systems: A practical approach to design, implementation, and management. 6th edn. Harlow, Essex, England: Pearson Education Limited.

13.1 Enhanced Entity-Relationship Model

13.2 Semistructured Databases (XML)

13.3 XML Data Manipulation (XPath and XQuery)

13.4 Transactions

13.5 Concurrency Control

13.6 Distributed Transactions

13.7 Distributed Concurrency Control

14 Distributed Systems

This section is taught by Dr Vasantam. The recommended reading for this topic is:

- Van Steen, M. and Tanenbaum, A.S. (2025). Distributed Systems. 4th edn. [online] Self-Published. Available at: <https://www.distributed-systems.net/index.php/books/ds4/> [Accessed 5 Oct. 2026].

14.1 Introduction to Distributed Systems

14.2 Middleware and RMI Technologies

14.3 Message-Oriented Middleware and Web Services

14.4 Replication Models

14.5 Concurrent Events and Consistency Control

14.6 Fault Tolerance

14.7 Load Distribution

Part IV

Programming Paradigms

15 Module Introduction

Programming Paradigms has the module code COMP2221 and is led by Dr Stuart James.

15.1 Assessment

This module is assessed through one piece of summative coursework and two final exams, these can be seen in Table 8. Additional formatives are also shown in *italics*.

Title	Opening	Due	Feedback	Professor	Weight
<i>Systems Programming</i>	<i>12th Oct</i>	<i>12th Nov</i>	<i>10th Dec</i>	<i>Dr Stuart James</i>	-
Systems Programming	23rd Oct	10th Dec	21st Jan	Dr Stuart James	50%
Object Oriented & Functional	May Exam	-	-	Ivkovic & Moses	50%

Table 8: Assessments for Programming Paradigms.

15.2 Timetable

This module is taught through two one-hour lectures and one two-hour computer class each week as seen below in Table 9.

Day	Time	Class Type	Location	Lecturer
Monday	11:00 - 12:00	Lecture	CLC202	Dr Stuart James
Monday	16:00 - 18:00	Computer Class	CB-LG001	n/a
Friday	14:00 - 15:00	Lecture	CLC202	Dr Stuart James

Table 9: Michaelmas Term Timetable for Programming Paradigms.

16 Systems Programming

This section is taught by Dr Stuart James. There is no official reading list for this module.

Definition 16.0.0.1: Systems programming involves the development of individual pieces of software that allow the entire system to function as a single unit.

16.1 Version Control and UNIX

16.1.1 Introduction to UNIX

16.2 The Basics of C

16.2.1 Basic File Setup

Definition 16.2.1.1: Pre-processor directives are lines which start with a `#` and are commands to the C pre-processor.

- If the directive is `#include`, and `<>` are used, then `/usr/include` is prioritised in UNIX by convention, whereas if `"` are used, then the current directory is prioritised.
- **Definition 16.2.1.1:** If the directive is `#define`, then this is a definition macro. These are used to provide definitions in code.

- **Definition 16.2.1.2:** You can also use directives for debugging without generating overheads, these are called conditionals. They are as shown in Table 10.

Conditional	Function
<code>#ifdef VAR</code>	Tests if VAR is <code>#defined</code> .
<code>#ifndef VAR</code>	Tests if VAR is not <code>#defined</code> .
<code>#else</code>	Runs if the original condition evaluates as false.
<code>#endif</code>	Ends the if testing condition.

Table 10: Different types of conditionals.

You can also put debugging code in the preamble, as seen in Algorithm 1.

```

#define MY_DEBUG // Define an identifier.

#ifdef MY_DEBUG
    assert(i > 0);
    printf("i is %d \n", i);
#endif

```

Algorithm 1: Debugging Code Example.

Definition 16.2.1.4: You can also use `#define` to define parameterised macros. These act like functions, and an example can be seen in Algorithm 2.

```

#define ADD(a,b) a + b

```

Algorithm 2: Parameterised Macro Example.

Definition 16.2.1.5: The main() function is the entry function for C programs. You can only have one of these.

16.2.2 Printf Parameters

Value	Resulting Format
<code>%d</code>	Signed decimal (int).
<code>%u</code>	Unsigned decimal.
<code>%o,%x</code>	Octal and hexadecimal.
<code>%l</code>	Long.
<code>%f</code>	Floating point, so <code>%1.2f</code> will give 3.14.
<code>%e</code>	Floating point (exponent form).
<code>%c,%s</code>	Character and string.

Table 11: First parameter of printf.

16.3 Control Flow and Functions

16.4 Data Types and Structures

16.5 Advanced Bash Scripting

16.6 Pointers

16.7 Memory Management

16.8 Debugging

16.9 Scope and Recursion

16.10 Libraries

16.11 stdlib

16.12 C++

17 Functional Programming

This section is taught by Dr Billy Moses Jr.. [Check for recommended reading.](#)

18 Object Oriented Programming

This section is taught by Dr Mladen Ivkovic. [Check for recommended reading.](#)

18.1 Java Fundamentals

18.2 OOP with Classes and Objects

18.3 Encapsulation

18.4 Inheritance

18.5 Abstraction with Abstract Classes

18.6 Abstraction with Interfaces

18.7 Polymorphism

18.8 Java Collections

18.9 Best Practises

Part V

Software Engineering

19 Module Introduction

Software Engineering has the module code COMP2281 and is led by Dr Craig Stewart.

19.1 Assessment

This module is assessed through a group project, the deadlines for this can be seen in Table 12. Additional formatives are also shown in *italics*.

Title	Opening	Due	Feedback	Professor	Weight
Peer Evaluation 1	16th Oct	19th Nov	10th Dec	Dr Craig Stewart	1.67%
Requirement Documentation	16th Oct	19th Nov	17th Dec	Dr Craig Stewart	15%
Peer Evaluation 2	16th Oct	11th Feb	11th Mar	Dr Craig Stewart	1.67%
Product Presentation	Week 16	-	18th Mar	Dr Craig Stewart	15%
Technical Report	5th Feb	18th Mar	6th May	Dr Craig Stewart	30%
Peer Evaluation 3	5th Feb	18th Mar	6th May	Dr Craig Stewart	1.67%
Git Commit	5th Feb	29th Apr	1st Jul	Dr Craig Stewart	5%
Project Handover	5th Feb	29th Apr	1st Jul	Dr Craig Stewart	10%
Reflective Experience Report	5th Feb	6th May	1st Jul	Dr Craig Stewart	20%

Table 12: Assessments for Software Engineering.

19.2 Timetable

This module is taught through two one-hour lectures and one two-hour computer class each week as seen below in Table 13.

Day	Time	Class Type	Location	Lecturer
Tuesday	09:00 - 10:00	Lecture	TLC033	Dr Nelly Bencomo
Wednesday	11:00 - 12:00	Lecture	TLC033	Dr Nelly Bencomo
Thursday	11:00 - 13:00	Practical	CC-0007	Dr Gagangeet Aujla

Table 13: Michaelmas Term Timetable for Software Engineering.

These lectures are taught jointly by Dr Craig Stewart and Dr Nelly Bencomo.

20 Project Management

21 Software Engineering

22 Legal, Social, Ethical, and Professional Issues

This recommended reading for this sub-topic is:

- The Software Engineering Code of Ethics and Professional Practice, ACM (2024).

22.1 Introduction to Software Engineering and Software Engineering Ethics

This lecture was taught by Dr Craig Stewart.

22.1.1 Introduction to Software Engineering

Definition 22.1.1.1: An ill-structured problem, as defined by (Simon, 1973), is: elusive, unbounded, not understandable, without end conditions. They require orientation and context, and have many stakeholders and many possible solutions.

Definition 22.1.1.2: A good problem solver needs the following:

- **Method:** Refers to a formal procedure for accomplishing a goal, typically independent of the tools used.
- **Tool:** An instrument or automated system for accomplishing something in a better way.
- **Procedure:** A combination of tools and techniques to produce a product.
- **Paradigm:** The philosophy or approach for building a project.

Definition 22.1.1.3: For this module, a stakeholder will generally fall into one of three roles:

- **Customer/Client:** The company, organisation, or person who pays for the software system.
- **Developer:** The company, organisation, or person who is building the software system.
- **User:** The person/people who will actually use the system.

22.1.2 Introduction to Ethics

Definition 22.1.2.1: The software engineering code of ethics is defined by Rost and Glass (2010) and summarised more thoroughly by the Association for Computing Machinery (2024).

23 Group Project

Part VI

Theory of Computation

24 Module Introduction

Theory of Computation has the module code COMP2181 and is led by Dr Barnaby Martin.

24.1 Assessment

This module is assessed through two summative benchtests, and one final exams, these can be seen in Table 14. Additional formatives are also shown in italics.

Title	Opening	Due	Feedback	Professor	Weight
Benchtest 1	Week 11	-	2nd Feb	Dr Dantchev and Dr Mertzios	12.5%
Benchtest 2	Week 20	-	-	Dr Dantchev and Dr Mertzios	12.5%
Final Exam	May Exam	-	-	Dr Barnaby Martin	75%

Table 14: Assessments for Theory of Computation.

24.2 Timetable

This module is taught through two one-hour lectures and one two-hour computer class each week as seen below in Table 15.

Day	Time	Class Type	Location	Lecturer
Tuesday	10:00 - 11:00	Lecture	CLC202	Dr Barnaby Martin
Wednesday	09:00 - 11:00	Problem Class	TLC116	-
Thursday	14:00 - 15:00	Lecture	CLC202	Dr Sukanya Pandey

Table 15: Michaelmas Term Timetable for Theory of Computation.

25 Models of Computation

This section is taught by Dr Barnaby Martin. [Check for recommended reading.](#)

25.1 Finite-State Automata and Regular Languages

25.1.1 Formal Definitions

Definition 25.1.1.1: A deterministic finite-state automaton is a five-tuple $(Q, \Sigma, \delta, q_0, F)$ where Q is the finite set of states, Σ is a finite alphabet, $\delta : Q \times \Sigma \rightarrow Q$ is the transition function, $q_0 \in Q$ is the start state, and $F \subseteq Q$ is the set of accept states.

Definition 25.1.1.2: A DFA M recognises the language L if $L = \{w | M \text{ accepts } w\}$.

Definition 25.1.1.3: A language is called a regular language if some DFA recognises it.

Definition 25.1.1.4: A regular expression R defines a language $L(R)$. We shall eventually prove that $RE \equiv DFA$. R is a regular expression over the alphabet Σ if R is:

1. a for some $a \in \Sigma$.
2. ϵ (the empty string).
3. $\emptyset$ (the empty set).

4. $(R_1 \cup R_2)$ where R_1 and R_2 are regular expressions.
5. $R_1 \cdot R_2$ where R_1 and R_2 are regular expressions.
6. R_1^* where R_1 is a regular expression.

The list of regular operations can be seen in Table 16 where the operations labelled with $^+$ are language theory specific.

Operation Name	Notation	Set Equivalent
Union	$A \cup B$	$\{x x \in A \text{ or } x \in B\}$
Intersection	$A \cap B$	$\{x x \in A \text{ and } x \in B\}$
Difference	$A \setminus B$	$\{x x \in A \text{ and } x \notin B\}$
Complement	$\overline{A}$	$\Sigma^* \setminus A$
Concatenation ⁺	$A \circ B$	$\{xy x \in A \text{ and } y \in B\}$
Star ⁺	A^*	$\{x_1 x_2 \dots x_k k \geq 0 \text{ and } x_i \in A \text{ for every } i, 1 \leq i \leq k\} \equiv \bigcup_{i=0}^{\infty} A^i$

Table 16: Regular Operations.

25.1.2 Combining Automata

Sometimes we may want to implement an operation on multiple languages. We can do this by combining their automata. If we have that L_1 is recognised by $M_1 = (Q_1, \Sigma, \delta_1, q_1, F_1)$ and L_2 is recognised by $M_2 = (Q_2, \Sigma, \delta_2, q_2, F_2)$; we want to combine M_1 and M_2 into a new automaton M that would recognise $L_1 \cup L_2$.

You would think that we could just run the automata sequentially, and accept if either M_1 or M_2 accept, but this does not work. This is because after M_1 has run on the input, the input is exhausted and there is no way to rewind it in order to run M_2 .

Instead, we can attempt to run M_1 and M_2 on the input in parallel. We do this by using pairs of states (one for each DFA). Formally, the components of the new automaton $M = (Q, \Sigma, \delta, q_0, F)$ are:

1. The set of possible states contains every possible cross-set pair of Q_1 and Q_2 .

$$Q = Q_1 \times Q_2$$

2. For every $s \in Q_1, u \in Q_2, a \in \Sigma$. The transition function contains every transition from a pair of states from Q being given any input from Σ .

$$\delta((s, u), a) = (\delta_1(s, a), \delta_2(u, a))$$

3. The start state is the pair containing both start states from M_1 and M_2 .

$$q_0 = (q_1, q_2)$$

4. The set of accept states is every possible cross-set pair of accept states from F_1 and F_2 .

$$F = \{(s, u) \in Q | s \in F_1 \text{ or } u \in F_2\}$$

25.1.3 Non-Deterministic Finite Automaton

If, instead you wish to work with concatenation, such that $w \in L_1 \circ L_2$ only if there are words x and y such that $w = xy$, $x \in L_1$ and $y \in L_2$. You need to run M_1 on a prefix of w and then run M_2 on the rest. This must be done non-deterministically and stops becoming a DFA as all accept states of M_1 have to map to the start state of M_2 .

Definition 25.1.3.1: A non-deterministic finite automaton is a five-tuple $(Q, \Sigma, \delta, q_0, F)$, where Q is the finite set of states, Σ is a finite alphabet, $\delta : Q \times (\Sigma \cup \{\epsilon\}) \rightarrow \mathcal{P}(Q)$ is the transition function, $q_0 \in Q$ is the start state, and $F \subseteq Q$ is the set of accept states.

Computing NFA

Before computing an NFA, an important convention to note is that a single step in a computation consists of first reading the next input symbol followed by taking any number of ϵ -transitions. This means that empty strings ϵ s can be freely added inside the actual word $\in \Sigma^*$, e.g., 010110 can be viewed as, say, $\epsilon 01\epsilon\epsilon 01\epsilon 01$.

Let $M = (Q, \Sigma, \delta, q_0, F)$ be an NFA and $w = w_1w_2\dots w_n$ be a word over $\Sigma \cup \{\epsilon\}$. M accepts w if there exists a sequence of states $r_0, r_1, r_2, \dots, r_n$ satisfying the following conditions:

1. $r_0 = q_0$.
2. $r_{i+1} \in \delta(r_i, w_{i+1})$ for every $i, 0 \leq i \leq n - 1$.
3. $r_n \in F$.

25.1.4 Converting NFA into DFA

The difference between a DFA and NFA is that a DFA has a unique computation (path) and either accepts or rejects, while a NFA generally has many computation paths and accepts if at least one computation accepts and only rejects if all the computations reject.

Reached slide 22/55 (powerpoint says slide 21).

25.2 Büchi Automata and Omega-Regular Languages

25.3 Linear-Time Logic and Translating Into Büchi Automata

26 Algorithms and Complexity

This section is taught by Dr Sukanya Pandey. The recommended reading for this topic is:

- Introduction to Algorithms (Cormen T.H., et al., 2009).

26.1 Asymptotic Complexity and Notation

Lecture one was just a recall of some of last year's Algorithms and Data Structures module.

26.2 Greedy Algorithms

26.3 Divide and Conquer Algorithms

26.4 Dynamic Programming Algorithms

26.5 Machinings in Graphs

Part VII

Appendix

27 Artificial Intelligence

27.1 Module Professors

Module Lead: Doctor Yang Long

- Focus: Scalable AI solutions for large-scale healthcare applications.
- Email: yang.long@durham.ac.uk
- Office Hours: Monday 15:30 - 16:30 in Term 1.

Lecturer: Professor Iain Stewart

- Focus: Computational complexity, graph/group theory, and theoretical aspects artificial intelligence.
- Email: i.a.stewart@durham.ac.uk
- Office Hours: Friday 14:00 - 15:00 in Term 2.

Lecturer: Doctor Shauna Concannon

- Focus: Computational linguistics, social interaction, and human-computer interaction.
- Email: shauna.j.concannon@durham.ac.uk
- Office Hours: Thursday 14:30 - 15:30 in Term 2.

27.2 References

27.3 List of Figures

27.4 Reading List

27.4.1 Machine Learning

Long, Y. and Durham University (2026). COMP2261.Machine.Learning.Study.and.Coursework.Guide-with.Navigation.pdf.

Deisenroth, M.P., Faisal, A.A. and Ong, C.S. (2020). Mathematics for Machine Learning. [online] Cambridge University Press. Available at: <https://mml-book.com/> [Accessed 3 Oct. 2026].

Hackeling, G. (2017). Mastering Machine Learning with scikit-learn. 2nd edn. Packt Publishing.

27.4.2 Ethics and Bias in AI

Various, will be added over time.

28 Data Science

28.1 Module Professors

Module Lead: Professor Thomas Erlebach

- Focus: Algorithms and networks.
- Email: thomas.erlebach@durham.ac.uk
- Office Hours: n/a.

Lecturer: Doctor Igor Razgon

- Focus: Algorithms and complexity, graph theory, and discrete mathematics.
- Email: igor.razgon@durham.ac.uk
- Office Hours: Friday 14:00 - 15:00.

Lecturer: Doctor George Alex Koulieris

- Focus: Image processing and computer graphics.
- Email: georgios.a.koulieris@durham.ac.uk
- Office Hours: Please email to make an appointment.

28.2 References

28.3 List of Figures

28.4 Reading List

28.4.1 Probability

Degroot, M.H. and Schervish, M.J. (2012). Probability and Statistics. Boston Mass.: Pearson.

28.4.2 Data Cleaning and Analytics

McKinney, W. (2022). Python for Data Analysis. O'Reilly Media, Inc.

29 Networks and Systems

29.1 Module Professors

Module Lead: Doctor Wanqing Tu

- Focus: Computer networks and wireless communications.
- Email: wanqing.tu@durham.ac.uk
- Office Hours: Thursday 11:00 - 12:00.

Lecturer: Doctor Thirupathaiah Vasantam

- Focus: Applied probability and performance modelling.

- Email: thirupathaiah.vasantam@durham.ac.uk
- Office Hours: Monday 14:00 - 15:00.

Lecturer: Doctor Carlton Shepherd

- Focus: Cyber security.
- Email: carlton.g.shepherd@durham.ac.uk
- Office Hours: Email to make an appointment.

29.2 References

29.3 List of Figures

29.4 Reading List

29.4.1 Cyber Security

Nakov, S., Stefanov, M. and Shideroff, M. (2018). Practical Cryptography for Developers. [online] sec. Encryption: Symmetric and Asymmetric Available at: <https://cryptobook.nakov.com/> [Accessed 9 Oct. 2026].

29.4.2 Databases

Connolly, T.M. and Begg, C.E. (2015). Database Systems: A practical approach to design, implementation, and management. 6th edn. Harlow, Essex, England: Pearson Education Limited.

29.4.3 Distributed Systems

Van Steen, M. and Tanenbaum, A.S. (2025). Distributed Systems. 4th edn. [online] Self-Published. Available at: <https://www.distributed-systems.net/index.php/books/ds4/> [Accessed 5 Oct. 2026].

30 Programming Paradigms

30.1 Module Professors

Module Lead: Doctor Stuart James

- Focus: Artificial intelligence, computer vision, human computer interaction, and digital humanities.
- Email: stuart.a.james@durham.ac.uk
- Office Hours: Email to make an appointment.

Lecturer: Doctor Billy Moses Jr.

- Focus: Algorithms and biologically inspired computing.
- Email: william.k.moses-jr@durham.ac.uk
- Office Hours: Email to make an appointment.

Lecturer: Doctor Mladen Ivkovic

- Focus: Performance prediction and optimisation.
- Email: mladen.ivkovic@durham.ac.uk
- Office Hours: Email to make an appointment.

30.2 References

30.3 List of Figures

31 Software Engineering

31.1 Module Professors

Module Lead: Doctor Craig Stewart

- Focus: Adaptive systems, HCI, and user modelling.
- Email: craig.d.stewart@durham.ac.uk
- Office Hours: Friday 11:00 - 12:00.

Lecturer: Doctor Nelly Bencomo

- Focus: Software engineering for autonomous, self-adaptive, AI and ML systems.
- Email: nelly.bencomo@durham.ac.uk
- Office Hours: Email to make an appointment.

Lecturer: Doctor Gagangeet Aujla

- Focus: Data-driven and data-intensive cyber-physical systems and digital twins.
- Email: gagangeet.s.aujla@durham.ac.uk
- Office Hours: Email to make an appointment.

Lecturer: Doctor Gokberk Kabacaoglu

- Focus: Machine learning and numerical applications.
- Email: gokberk.kabacaoglu@durham.ac.uk
- Office Hours: Thursday 14:00 - 15:00.

Lecturer: Doctor Ryan Crosby

- Focus: The assessment of novice programmers.
- Email: ryan.crosby@durham.ac.uk
- Office Hours: Email to make an appointment.

Lecturer: Doctor Sukanya Pandey

- Focus: Graph theory.
- Email: sukanya.pandey@durham.ac.uk
- Office Hours: Email to make an appointment.

31.2 References

31.2.1 Legal, Social, Ethical, and Professional Issues

Simon, H.A. (1973). The Structure of Ill Structured Problems. Artificial Intelligence, 4(3-4).

Rost, J. and Glass, R.L. (2010). The Dark Side of Software Engineering: Evil on Computing Projects. Washington, DC: IEEE Computer Society.

31.3 List of Figures

31.4 Reading List

31.4.1 Legal, Social, Ethical, and Professional Issues

Association for Computing Machinery (2024). The Software Engineering Code of Ethics and Professional Practice. [online] acm.org. Available at: <https://www.acm.org/code-of-ethics/software-engineering-code> [Accessed 7 Oct. 2026].

32 Theory of Computation

32.1 Module Professors

Module Lead: Doctor Barnaby Martin

- Focus: Complexity theory, finite model theory, and the links between logic and complexity.
- Email: barnaby.d.martin@durham.ac.uk
- Office Hours: Flexible, contact to arrange.

Lecturer: Doctor Sukanya Pandey

- Focus: Graph theory.
- Email: sukanya.pandey@durham.ac.uk
- Office Hours: Flexible, contact to arrange.

Lecturer: Doctor Stefan Dantchev

- Focus: Combinatorics.
- Email: s.s.dantchev@durham.ac.uk
- Office Hours: **To be confirmed.**

Lecturer: Doctor George Mertzios

- Focus: Algorithms and networks.
- Email: george.mertzios@durham.ac.uk
- Office Hours: **To be confirmed.**

32.2 References

32.3 List of Figures

32.4 Recommended Reading

32.4.1 Algorithms and Complexity

Cormen, T.H., Leiserson, C.E., Rivest, R.L. and Stein, C. (2009). Introduction to Algorithms. 3rd edn. MIT Press.